Closure-Based Syntax for Contracts

Contents

1   Introduction

The attribute-based syntax for contracts is limiting and steps on the shared space between C and C++. This paper explores an alternative syntax that should offer an easier extension path.

This paper proposes almost the same semantics as [P2388R2].

The only significant change from [P2388R2] is the semantics of effect elision - this paper specifies it as all-or-nothing, per *correctness-annotation.*

The main non-syntactic difference between the two papers is the manner of spelling future extensions.

Due to the way this paper models annotations, it may leave fewer things undefined compared to [P2388R2], despite the fact that that it does not propose explicit closures yet.

Note: this paper is an exploration. The authors do not strongly object to the attributes-based syntax from the P2388 series; but the syntax does seem to present certain challenges that this paper tries to address.

Note: WG14 has communicated that their vendors don't have a blocking problem with the attribute-like syntax, though they have reservations; in addition, WG21 members have expressed difficulties with teaching the *: means it's not an attribute* intricacies.

1.1   On Extensions and Viability

The authors believe that **any MVP must clearly show plausible syntax for all known extensions**. This does not mean *propose*. It means *show*. Specifying precise semantics for the entire extension space is not in the spirit of a *minimum viable product*, but *viability* implies that all desired features can at some point be supported. This means there must be syntax, so syntax we show.

2   Changelog

2.1   R0 to R1

— general clarifications and better choice of words, fixed loads of typos.
— The grammar specification separated into *precondition*s, *postcondition*s, and *assertion*s, to enable more precision in specification.

— the return-value parameter to a *postcondition* was made optional.
— The default capture case for member functions was changed to `[&, this]`. This fixes an oversight.
— We clarify that the closure was always supposed to be as-if `mutable` in the by-value capturing extensions.
— Greatly expanded examples, and clarified which ones are extensions and which ones are proposed here.
— Removed "lambda" from the paper and consistently use "closure", as the mechanism for capturing values at function entry relies on lambda closures and their syntax, but we don't rely on the lambda body syntax.
— Split the "future extensions" sketch chapter into subclauses highlighting individual separable extensions.
— Highlight the different models of side-effect elision (see "An alternative model"), and what we give up by using a looser one than proposed here.
— Changed the destructuring of the return value example into an actually interesting one.
— Changed the location of the *attribute-specifier-seq$_{opt}$* in the extension allowing attributes appertaining to contract annotations.
— Added a section on testing and fuzzing contract specifications themselves.
— Clarified the relationship to abbreviated lambdas
— Added the "capture design space" section.
— Added section on considered and rejected ideas
— Clarified that contract checks are ODR-used even in *nocheck* mode.

3 Proposal

3.1 Example

We introduce three context-sensitive keywords: `pre`, `post`, and `assert`. `pre` and `post` are only keywords in the top level context of a function declarator.

`pre` and `post` can appear in function declarations after the optional trailing `requires` clause.

Example:

```
auto plus(auto const x, auto const y) -> decltype(x + y)
    pre { x > 0 }
    pre { p1(x) && p2(x) } // compound check
    // r is as-if auto&&
    post (r) { r == (x + y) }
{
    assert { x > 0 }; // this is currently "valid" syntax,
                      // but we should reclaim it.
    auto cx = x;
    return cx += y;
}
```

One may note that this is strikingly similar to the syntax proposed in [N1962], way back in 2006. Our thanks to Andrzej Krzemiński for digging this up.

3.2 Proposed syntax

Let's take a look at the generic syntax of a *correctness-annotation* (to use the term from [P2388R2]):

correctness-specifier:
  precondition
  postcondition
  assertion

precondition:
  `pre` *lambda-introducer$_{opt}$* *correctness-specifier-body*

postcondition:
    `post` *lambda-introducer$_{opt}$ return-value-decl$_{opt}$ correctness-specifier-body*

assertion:
    `assert` *lambda-introducer$_{opt}$ correctness-specifier-body*

return-value-decl:
  (*identifier*)

correctness-specifier-body:
  { *conditional-expression* }

For the MVP, the *lambda-introducer* **is required to be omitted**.

If the *lambda-introducer* is omitted, the *correctness-specifier-body* behaves as-if the *lambda-introducer* was `[&]` for free functions and `[&, this]` for member functions.

In a *postcondition*, the *return-value-decl*, if present, introduces the name for the prvalue or the glvalue result object of the function. This identifier is valid within the *correctness-specifier-body*. This is to support annotating `void`-returning functions, as well as *postcondition*s that don't operate on the return value.

All closures behave as-if their associated lambda body was declared `mutable`. This makes no difference to `[&]`, but it does make a difference for by-value capture extensions.

3.2.1 MVP Restrictions

Naming a non-`const` value parameter in a *postcondition* is ill-formed for now. This can be lifted by allowing copy-capture later, when we allow the *lambda-introducer* to appear. This is to both prevent **referencing moved-from objects**, and to **allow the calling code to reason** about the properties of the result object, such as in the example:

```cpp
int min(int x, int y)
    post (r) { r <= x && r <= y }; // error, x and y are not const

int min(int const x, int const y)
    post (r) { r <= x && r <= y }; // ok

int min(int x, int y)
    post [x, y] (r) { r <= x && r <= y }; // OK (extension, explicit copy)
```

The closure definition works with this - the function arguments are captured by reference, which happens to be reference-to-`const`, given that they are `const`, which gives the exact semantics of [P2388R2].

4 Semantics

We specify the future in a somewhat more general manner than strictly required for the MVP, to indicate the inner workings of the future extensions.

4.1 Evaluation order

This section describes the order of evaluation *if contract checking is enabled*. If it's disabled, there is no evaluation, but the contract check bodies are still ODR-used.

4.1.1 Assertions

Any *assertion* is executed as if it was an immediately-invoked lambda expression.

4.1.2 *pre* and *postcondition*s

We need to make preceding *precondition*s protect both the *lambda-introducer* and the *correctness-specifier-body* of any subsequent *correctness-specifier*.

Therefore, *precondition*s are first execute the *correctness-specifier-introducer* (if any), and then immediately their *correctness-specifier-body*.

*postcondition*s are evaluated in two parts; their *correctness-specifier-introducer* is evaluated in-sequence along with *precondition*s; and their *bodies* are, evaluated after the function exits.

If a *precondition B* follows a *precondition A* in a function's declaration, then no part of *B* shall be executed before *A* has been proven;

If a *postcondition P* follows a *precondition A* in a function's declaration, then not even *P*'s *correctness-specifier-introducer* shall be executed before *A* is proven. This is to protect initialization from out-of-contract behavior.

No postcondition closure is executed before all *precondition*s are proven.

This means that the following execution orders are all OK:

— *A*, *B*, *P*
— *A*, *A*, *B*, *B*, *P*
— *A*, *B*, *A*, *B*, *P*
— *A*, *B*, *P*
— (prove *A* at compile time), *B*, *P*
— (inherit proof of *A* from caller precondition), *B*, *P*

4.2 *postcondition* reference-capture limitations in the MVP

Capturing function parameters by mutable-reference in *postcondition*s may cause difficulties for static analysis, as some expressions containing these will require interprocedural/inter-TU analysis, which may be beyond the capabilities of a compiler. Dedicated static analysis tools should still be able to handle these, however. [P2388R2] forbids mutating function arguments.

Example (courtesy of Tomasz Kamiński):

```
int pickRandom(int beg, int end)
    post [&] (r) {
        ret >= beg &&
        ret <= end
    };
```

Given that we don't know the function body, and we could have changed `beg` and `end`, this conveys no information for static analysis (you'd have to mark `beg` and `end const`).

We therefore have a choice of how to start out with this proposal:

— forbid capturing parameters by mutable reference
— forbid capturing parameters by reference altogether
— do nothing and just expect degraded static analysis performance (capture-by-mutable-reference is not a problem for runtime checking)

The stated goal of feature-bijection with [P2388R2] for this paper says we should forbid reference-capture for parameters in post-conditions.

4.3 Side-effect elision

This MVP defines that for the purposes of optimization, the compiler is allowed to either execute, or not, entire correctness specifiers, together with their closures. Subexpression elimination is only permitted under the (stricter) as-if rule.

This is because, while it should not be lippincott-discernible to the program whether a specifier was actually executed, this might only actually be true if the specifier gets to clean up after itself. In other words, the sum of the parts is assumed "pure", the parts are not.

4.3.1   An alternative model (from P2388)

P2388 has a looser model, which the authors of this paper do not oppose. In that model, the compiler is allowed to elide any and all side-effects indiscriminately, from any subexpression.

That model effectively disallows lock/unlock pairs inside the body of a check, but is unclear about lock/unlock pairs inside closure initialization, which P2388 does not speak of.

5   Future Extensions (not a proposal)

As noted in the introduction, the reason to prefer this paper to an attribute-based syntax is the way evolution is handled.

This section shows the various ways future extensions could look like under this proposal. These are not hypothetical, and will be proposed immediately after this MVP is accepted.

5.1   Explicit Captures

The main task of this proposal is enabling explicit parameter captures in the future. The main consumer of that are *postcondition*s.

```
auto plus(auto x, auto y) -> decltype(x + y)
    post [x, y] (r) {
        // capture x, y by value, *explicitly*, at point of call
        r == (x + y)
    }
{
    return x += y;
}
```

Modeling using lambda-captures allows us to explain why post-conditions can't refer to rvalue-reference arguments, and all the other possible implementation-limitations as well.

5.1.1   Capturing view contents by value

Explicit captures allow capturing views by value for later checking. No "magic" *oldof* implementation can do this - but closures do this easily.

```
template <typename T, size_t n>
void sort_contiguous(std::span<T, n> elems)
  post [on_entry = std::vector(elems), elems] { is_permutation(on_entry, elems) }
  post [elems] { std::is_sorted(elems.begin(), elems.end()) }
{
  std::sort(elems.begin(), elems.end());
}
```

5.1.2   Checking whether a call didn't exceed its time budget

A yet-unserved use-case is checking whether a realtime function actually runs in the time promised; this syntax makes it easy:

```
int runs_in_under_10us()
    post [start=gettime()] { gettime() - start <= 10us };
```

The authors thank Lisa Lippincott for this wonderful idea during a conversation in Aspen a few years ago.

5.1.3 Checking a call didn't leak memory

With a tracking allocator, we can check we didn't leak any memory:

```cpp
int does_not_leak(allocator auto alloc)
    post [usage=alloc.usage(), &alloc] { usage == alloc.usage() }
{
  // do stuff that should not leak
  return 0;
}
```

5.1.4 Grabbing only the interesting part of an input

We can optimize contracts by "remembering" just the required properties of an input:

```cpp
void append(auto& container, auto&& item)
  post [s=container.size()] { container.size() == s+1 }
{
  container.push_back(std::forward<decltype(item)>(item));
}
```

5.1.5 Mutation semantics checking

Sometimes we want to check that two operations are equivalent for the given inputs, because the algorithm is taking advantage of that.

```cpp
auto cat(auto x, auto y)
  pre [cx=x] { (x+=y) == x+y }
{
  return std::move(x += y);
}
```

5.2 Destructuring the return value

There have been rumours of a proposal for destructuring in function arguments. When the language gets that, we can just inherit that directly. A teaser that we *could* just adopt directly:

```cpp
auto returns_triple()
    post ([a, b, c]) { c > 0 }
{
    struct __private { int __a; int __b; int __c; };
    return __private{1, 2, 3};
}
```

5.3 Attributes appertaining to contract annotations

The syntax allows for attributes on annotations. We could use those for vendor-defined control of execution.

```cpp
int f(int * n)
    pre                 {n != nullptr}
    pre [[acme::audit]] {*n >= 0};
```

An alternative (although the authors are not completely sure this does not conflict with the *attribute-specifier-seq* that appertains to the function type):

```
int f(int * n)
    [[acme::audit]] pre {*n >= 0};
```

This one courtesy of Andrzej Krzemiński.

5.4 "trust" annotations (new)

Comparison from [P2388R2]/8:

Extension of this proposal	[P2388R2]
```int f(int* p) pre new {*p > 0};```	```// after ; at end int f(int* p) [[pre: *p > 0; new]] ;```
```int f(int* p) pre new("call @me") {*p > 0};```	```// after ; at end int f(int* p) [[pre: *p > 0; new("call @me")]] ;```

5.5 "cost" annotations

In the same way, we can add convenient cost annotations:

Extension of this proposal	[P2388R2]
```void sort(auto first, auto last) post audit("really expensive") [s=vector(first, last)] { is_permutation({first, last}, s); }```	```void sort(auto first, auto last) [[post r: r > 0: audit("really expensive")]];```
	```void sort(auto first, auto last) [[post audit("really expensive") r: r > 0:]];```

5.6 Multithreaded usage / locking

Issue courtesy of Aaron Ballman:

A potential issue with P2388R2 that is carried over into D2461R0 is with side effect operations. Given that they're unspecified, does this mean there's no safe way to write a portable contract which accesses an object shared between threads? e.g., multithreaded program where a function is passed a mutex and a pointer to a shared object; can the contract lock the mutex, access the pointee, then unlock the mutex?

With closure-based semantics, we can avoid this:

```
void frobnicate_concurrently(auto&& x)
    // closures-are-a-future-extension.disclaimer
```

```
    pre [g=std::lock_guard(x)] { is_uniquely_owned(x); };
```

In this MVP, we allow the compiler to *assume there are no side-effects* to an expression for the purposes of optimisation, *but they can either all be omitted, or none may*, for a given statement, including the closure.

We therefore have a plausible RAII-based metaphor that people already understand.

5.7 Testing and fuzzing *precondition*s and *postcondition*s

It's important to be able to test the actual precondition and postcondition specification independent of the function they are guarding.

This is due to

— enabling contracts should ideally not introduce *additional* bugs.
— contracts necessarily *narrow* from the *wide* contract specification; therefore the contract specification of the *check* is wider than the contract of the *function*.
— Contracts need to "expect the unexpected". This makes them excellent candidates for fuzzing.

We could introduce a pair of reflection traits for functions:

```
// unchecked precondition: y != 0
int f(int x, int y)
  pre { x / y >= 0 }
  pre [x] { x %= 3 != 0 }
  post (r) { r == x / y }
;

// all *precondition*s
std::preconditions<f>(/*x*/3, /*y*/2); // false, x%=3 == 0
// by-index
std::preconditions<f>[0](3, 2); // true, 3 / 2 == 1 , >= 0
std::preconditions<f>[1](3, 2); // false, 3 / 2 == 1 , >= 0

// all *postcondition*s
std::postconditions<f>(/*x*/3, /*y*/2)(/*r*/1); // true, 1 == 3/2
std::postconditions<f>[0](3, 2)(1); // true, 1 == 3/2
std::postconditions<f>[0](4, 2)(1); // false, 1 != 4/2

// fuzzer:
std::postconditions<f>[0](3, 0)(1); // segfault, bugreport
```

*postcondition*s *have* to be fuzzable independently of their functions. If the algorithm is wrong (the very condition the postcondition is there to detect), the postcondition should absolutely not segfault, but instead report `false`.

Traits such as above absolutely must be independently executable.

We could also introduce names:

```
int g(int x)
  pre("nonnegativity") { x >= 0 }
  post("nonnegativity") (r) { r >= 0 }
;

std::preconditions<g>["nonnegativity"](1); // true
std::postconditions<g>["nonnegativity"](1); // true
```

We could propose the same for attribute-based syntax, but the closure + argument model has clear

semantics, and with the other alternative we'd have to additionally specify how it works. Support for names takes some syntactic space, which is already scarce in the attribute-based syntax.

5.8 Summary

— We get improvements in lambda-capture grammar "for free". Once lambda-introducers get destructuring support, so do contracts, instead of inventing yet another minilanguage.
— We don't have to re-specify anything regarding pack expansions, etc; lambda-introducers get us that, too.
— We can check time/environment-based contracts (see example below).
— It's consistent with the rest of the language, instead of inventing a yet-another minilanguage.

6 Comparison tables with attribute-based syntax

This section explores future extensions as envisaged by [P2388R2] and previous papers. The comparison is also is explored in [P2487R0].

6.1 Referencing function arguments in *postcondition*s

There are issues with arguments that change value during function evaluation and *postcondition*s. They are described in [P2388R2]/6.4 and 8.1. [P2388R2] side-steps this issue by attempting to prevent referencing modified arguments, requiring that referenced arguments should be `const`-qualified (in definitions).

The ideas using the [P2388R2] syntax look like this (all from [P2388R2]/8.1):

```
// Extension of this proposal
int f(int& i, array<int, 8>& arr)
  post [i] (r) { r >= i }
  post [old_7=arr[7]] (r)
        { r >= old_7 }
```

```
// p2388r2 1)
int f(int& i, array<int, 8>& arr)
  [[post r, old_i = i: r >= old_i]]
  [[post r, old_7 = arr[7]: r >= old_7]];
```

```
// p2388r2 3)
int f(int& i, array<int, 8>& arr)
  [[post r: r >= oldof(i)]]
  [[post r: r >= oldof(arr[7])]];
```

```
// p2388r2 2)
int f(int& i, array<int, 8>& arr)
  [[post r: r >= oldof(i)]]
  [[post r: r >= oldof(arr[7])]];
```

Table 4: Another oldof example:

Extension of this proposal	P2388R2
``template<class ForwardIt, class T>
ForwardIt find(ForwardIt first,
               ForwardIt last,
               const T& value)
  post [first] (r)
       { distance(first, r) >= 0u }
  post [&last] (r)
       { distance(r, last) >= 0u }
{
  for (; first != last; ++first) {
    if (*first == value) {
      return first;
    }
  }
  return last;
}``` | ```template<class ForwardIt, class T>
ForwardIt find(ForwardIt first,
               ForwardIt last,
               const T& value)
  [[post r: distance(oldof(first), r) >= 0u]]

  [[post r: distance(r, last) >= 0u]]

{
  for (; first != last; ++first) {
    if (*first == value) {
      return first;
    }
  }
  return last;
}``` |

6.2 Introducing the return variable

Extension of this proposal	P2388R2
``int f(int* i, array<int, 8>& arr)
post [&i] (r) { r >= i };```	```int f(int& i, array<int, 8>& arr)
[[post r: r >= i]];```
	``// alternative
int f(int& i, array<int, 8>& arr)
[[post(r): r >= 0]]```

6.3 *precondition*s and assertions that need copies

Table 6: [P2388R2] has no answer for *precondition*s that need to mutate a copy:

Extension of this proposal	[P2388R2] Does not work
``int f(forward_iterator auto first,
forward_iterator auto last)
pre { first != last }
pre [first] { std::advance(first, 1),
first != last };```	```int f(forward_iterator auto first,
forward_iterator auto last)
[[pre: first != last]] // ok
[[pre: std::advance(first, 1), // error
first != last]]];```

6.4 *postcondition*s that need destructuring [when lambda-captures get it]

Table 7: Functions could conceivably have destructure-only APIs:

Extension of this proposal	[P2388R2]
```auto returns_triple() post ([x, y, z]) { x > y && y > z } ;```	```auto returns_triple() [[post [x, y, z]: x > y && y > z]];```

This syntax kind-of works, but is not proposed, and there is nowhere to specify the binding type (reference or copy?) in either of the possibilities.

The non-attribute based syntax has one advantage here, which is a clear place for the return value. If we put an additional closure initialization somewhere before the colon, we end up in visual ambiguity of what `[]` means.

6.5 Summary

— The closure-based syntax makes it obvious when values are captured, and even hints at an implementation - just put the closures on the stack before the function arguments.
— It doesn't invent another language for capturing values, which means the syntax will grow together with lambda captures.
— It makes it **obvious how to do stateful *postcondition*s** that check before/after: the closure runs with `pre`, the body runs after return. This is far from obvious with the [P2388R2] syntax.

7 Mutation and Static Analyzers

Static analyzers should be able to handle limited mutation in order to analyze C++, and many contracts that describe function behaviour will require some mutation of a copy. Allowing copies to be made is therefore immensely useful in a contract facility.

We have assurances from at least some analyzer vendors they see no issue with allowing copies and mutation in contract annotations in the future.

7.1 Capture design space

There is design space in what kind of captures we allow of parameters of different kinds.

The case space spans between condition (`pre`, `post`) by parameter kind (value, lvalue, rvalue) by qualifier (non-const, `const`) by kind-of-capture (reference, value), which gives us $2 * 3 * 2 * 2 = 24$ cases.

The example is

```
auto f(param x)
    pre cap { x }
    post cap { x }
;
```

#	param	cap	cond	MVP	Ext
1	T	[&]	pre	y	y
2	T	[&]	post	(!)	(!)
3	T	[=]	pre		y
4	T	[=]	post		y

#	param	cap	cond	MVP	Ext
5	T const	[&]	pre	y	y
6	T const	[&]	post	y	y
7	T const	[=]	pre		y
8	T const	[=]	post		y
9	T &	[&]	pre	y	y
10	T &	[&]	post	y	y
11	T &	[=]	pre		y
12	T &	[=]	post		y
13	T const &	[&]	pre	y	y
14	T const &	[&]	post	y	y
15	T const &	[=]	pre		y
16	T const &	[=]	post		y
17	T &&	[&]	pre	y	y
18	T &&	[&]	post	y	y
19	T &&	[=]	pre		y
20	T &&	[=]	post		y
21	T const &&	[&]	pre	y	y
22	T const &&	[&]	post	y	y
23	T const &&	[=]	pre		y
24	T const &&	[=]	post		y

— **MVP**: this paper.
— **Ext**: this paper plus minimal "allow captures" extension.
— **(!)**: Capturing value arguments by reference in *postcondition*s gives no information to the caller, therefore it's not a well-behaved contract and we should diagnose (2).
— All the [=]-captures are missing for the MVP due to the lack of explicit captures, which means there is no way to spell that.

The reader is encouraged to specify their own "acceptability mask" if they want to limit this further, and propose it in discussion. However, the authors don't really see how any further restriction from Ext couldn't be circumvented by the user, and (2) is protection from Murphy, not Machiavelli.

8 C-compatibility

C and C++ implementations often share a set of system headers, and there will naturally be a desire to add contracts to entities in those headers.

One of the motivating reasons behind the attribute-like syntax in [P2388R2] is that a C compiler can be reasonably updated to ignore the contracts unless/until C gets Contracts as well. It's worth noting that the proposed syntax in [P2388R2] is still ill-formed for attributes, and a properly conforming C compiler that has not been updated to handle (ignore) the contracts would still issue diagnostics.

There is some debate as to whether it'd be a *good* thing if a C compiler were to still accept code that has Contracts in it when the C compiler is unaware of Contracts, and it has been noted that some implementations may simply consume all tokens in an unrecognized attribute until reaching the closing]], regardless of whether the internal structure of the attribute is properly conforming.

The syntax proposed in *this* paper, however, cannot be ignored by a C compiler that is unaware of Contracts - it is unarguably ill-formed C code.

This syntax lends itself easily to conditional compilation, especially with a feature-test macro:

```
int my_func(int x)
#if __cpp_contracts /* Perhaps just __contracts to allow C to easily opt-in? */
```

```
  pre { x > 0; }
#endif /* __cpp_contracts */
{
  /* ... */
}
```

This is not a motivating difference from [P2388R2] - conditional compilation can just as easily be used to guard Contracts there; the main difference in C-compatibility between these two proposals is that [P2388R2] has a greater potential of a Contracts-unaware C compiler ignoring any contracts without a meaningful diagnostic or programmer opt-in.

9 Considered and rejected ideas

This section is a log of things the authors considered and rejected, and for which reason, so that we need not revisit prior discussions.

9.1 Abbreviated lambdas

The *postcondition* syntax naturally looks like a shorthand lambda:

```
post [ closure ] ( identifier ) { conditional-expression }
```

This topic was explored in [P0573R2] by Barry Revzin, who proposed this syntax as point 2.3. (alternative syntax). Notably, this was not the main proposed syntax, which was

```
post [ closure ] ( identifier ) => conditional-expression
```

In function declarations, this syntax ends up looking like

```
int f(auto x) pre => x > 0 pre => 42 % x == 0 post [x](r) => x * r == 42 {
  return 42 / x ;
}
```

It's *parsable*, but it doesn't scan well to the human eye. Contrast with the proposed

```
int f(auto x) pre { x > 0 } pre { 42 % x == 0 } post [x](r) { x * r == 42 } {
  return 42 / x ;
}
```

It has the same number of characters, but the visual terminator that } provides makes everything a lot more readable.

Reasons to reject:

— We do not want to wait for EWG to express a stance on abbreviated lambdas
— We want the closure to be optional, which is a departure
— The proposal does not consider making the closure `mutable` by default, which we want. The author said this was not considered at all and may be revisited in a future revision, but that does not exist at present.

9.2 Semicolon separators in the body (a-la requires blocks)

The authors also considered separating conditions with semicolons, similarly to `requires` blocks.

```
int f(auto x)
  requires {
    {x % 2} -> integral;
    {as_signed(x)} -> convertible_to<int>;
  }
```

```
  pre {
    x % 2 == 0; // short-circuits
    as_signed(x) == x;
  }
;
```

While it makes sense with `requires` blocks, the behavior of `;` in *pre-* and *postcondition*s is identical to the `&&`-operator, because we want short-circuiting. In addition, the lack of a `return` keyword in the body makes the individual conditions seem disembodied.

Requiring chaining using `&&` solves this comprehension issue – it behaves the way it looks, and the lack of a semicolon in the body hints at the fact that the thing between `{}` is an expression and not a statement.

In other words, we say what we mean, and we disallow the syntax that doesn't mean what it looks like.

10 Proposed Wording

TODO. Writing it will be an exercise, and the authors want to see if there is any enthusiasm for this at all before spending the time.

11 Acknowledgements

— *The entire WG21.* This is a huge effort, and since contracts were pulled from C++20, the group has been showing an extraordinary level of determination to get to consensus.
— *Andrzej Krzemieński*, who has been a steadfast integrator of opinion in P2338 - the MVP paper sequence. I've helped a bit, but he's been extraordinary, and also dug up more prior art and contributed examples. *I thrice presented him co-authorship, which he did thrice refuse.*
— *Tom Honermann*, who saw the interplay with function try-blocks.
— *Phil Nash*, for quite a few insightful comments, and the function-parameter syntax for return values
— *Peter Brett*, for encouraging me to drop the complex sequence of ;-separated conditions and stick to a single condition (subconditions separated by `&&`).
— The *BSI* for reviewing this paper early.
— *Lisa Lippincott*, for her study of stateful function contracts and all the hours she's spent explaining the point and their shape to Gašper.
— *Tomasz Kamiński*, for *also* pointing out the function parameter syntax for return values, and reminding the authors that reference-captures for non-`const` parameters render *postcondition*s less useful for static analysis in the absence of the function body.
— *Ville Voutilainen*, for always connecting all of the weird bits of impact everything has on everything else.
— *Bengt Gustaffson*, for an amazingly long and thorough review of the R0.

12 References

[N1962] L. Crowl, T. Ottosen. 2006-02-25. Proposal to add Contract Programming to C++ (revision 4).
https://wg21.link/n1962

[P0573R2] Barry Revzin, Tomasz Kamiński. 2017-10-08. Abbreviated Lambdas for Fun and Profit.
https://wg21.link/p0573r2

[P2388R2] Andrzej Krzemieński, Gašper Ažman. 2021-09-10. Minimum Contract Support: either Ignore or Check_and_abort.
https://wg21.link/p2388r2

[P2487R0] Andrzej Krzemieński. Attribute-like syntax for contract annotations.
https://isocpp.org/files/papers/P2487R0.html